

Update on Ara

12/05/2021

Matteo Perotti

Matheus Cavalcante

Nils Wistoff

Professor Luca Benini

Integrated Systems Laboratory

ETH Zürich

Summary

- **Toolchain**
 - From GCC to LLVM
 - CI pipeline
- **Hardware**
 - Transition to RVV 0.10
 - Implementing reductions
- **Software**
 - New benchmarks

Toolchain - From GCC to LLVM

Motivation:

“Current focus for V compiler work is on LLVM.” [Jim Wilson]




- Scalar riscv-tests
- Vector riscv-tests
- Vector programs

Toolchain - From GCC to LLVM

Motivation:

“Current focus for V compiler work is on LLVM.” [Jim Wilson]



- Scalar riscv-tests



- Vector riscv-tests
- Vector programs

Toolchain - From GCC to LLVM



- RVV 0.10
- V-intrinsics support
- Implies:
 - SPIKE RVV 0.10
 - Ara updated to RVV 0.10*

Hardware - Reductions

vs1

e0	e1	e2	e3	e4	e5	e6	e7	e8	e9	e10	e11
----	----	----	----	----	----	----	----	----	----	-----	-----

vs2

s0	-	-	-	-	-	-	-	-	-	-	-
----	---	---	---	---	---	---	---	---	---	---	---

**vd**

r0	-	-	-	-	-	-	-	-	-	-	-
----	---	---	---	---	---	---	---	---	---	---	---

Hardware - Reductions

Example: vredsum vd, vs1, vs2

vs1

e0	e1	e2	e3	e4	e5	e6	e7	e8	e9	e10	e11
----	----	----	----	----	----	----	----	----	----	-----	-----

vs2

s0	-	-	-	-	-	-	-	-	-	-	-
----	---	---	---	---	---	---	---	---	---	---	---



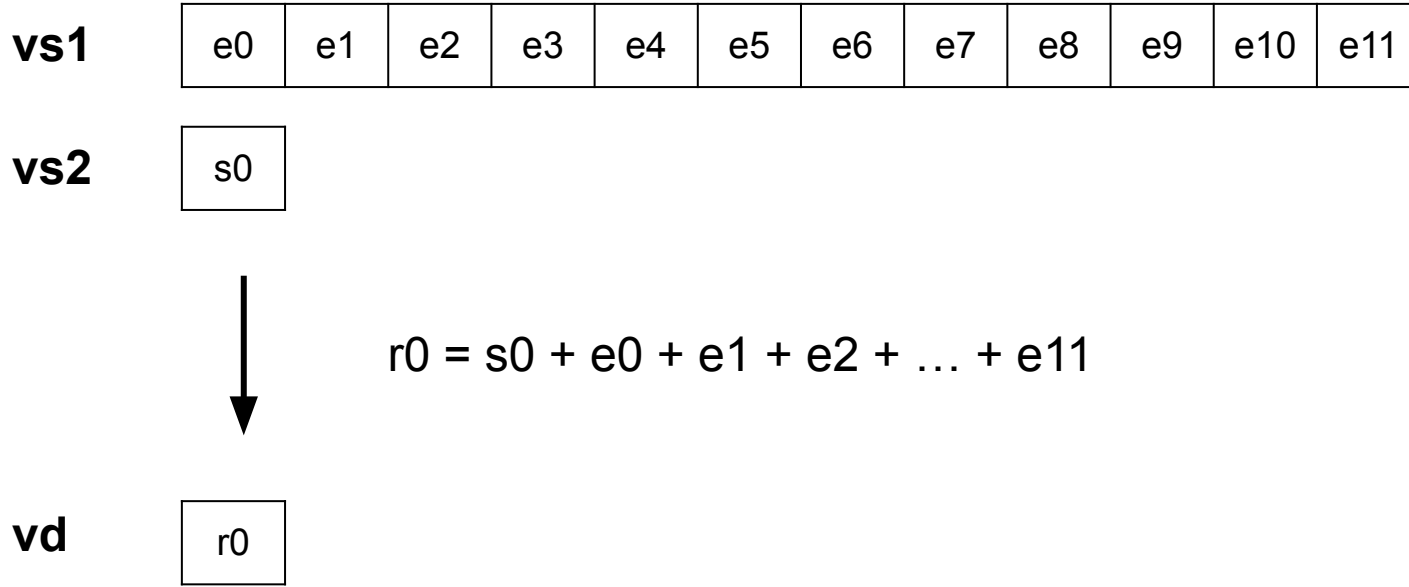
$$r0 = s0 + e0 + e1 + e2 + \dots + e11$$

vd

r0	-	-	-	-	-	-	-	-	-	-	-
----	---	---	---	---	---	---	---	---	---	---	---

Hardware - Reductions

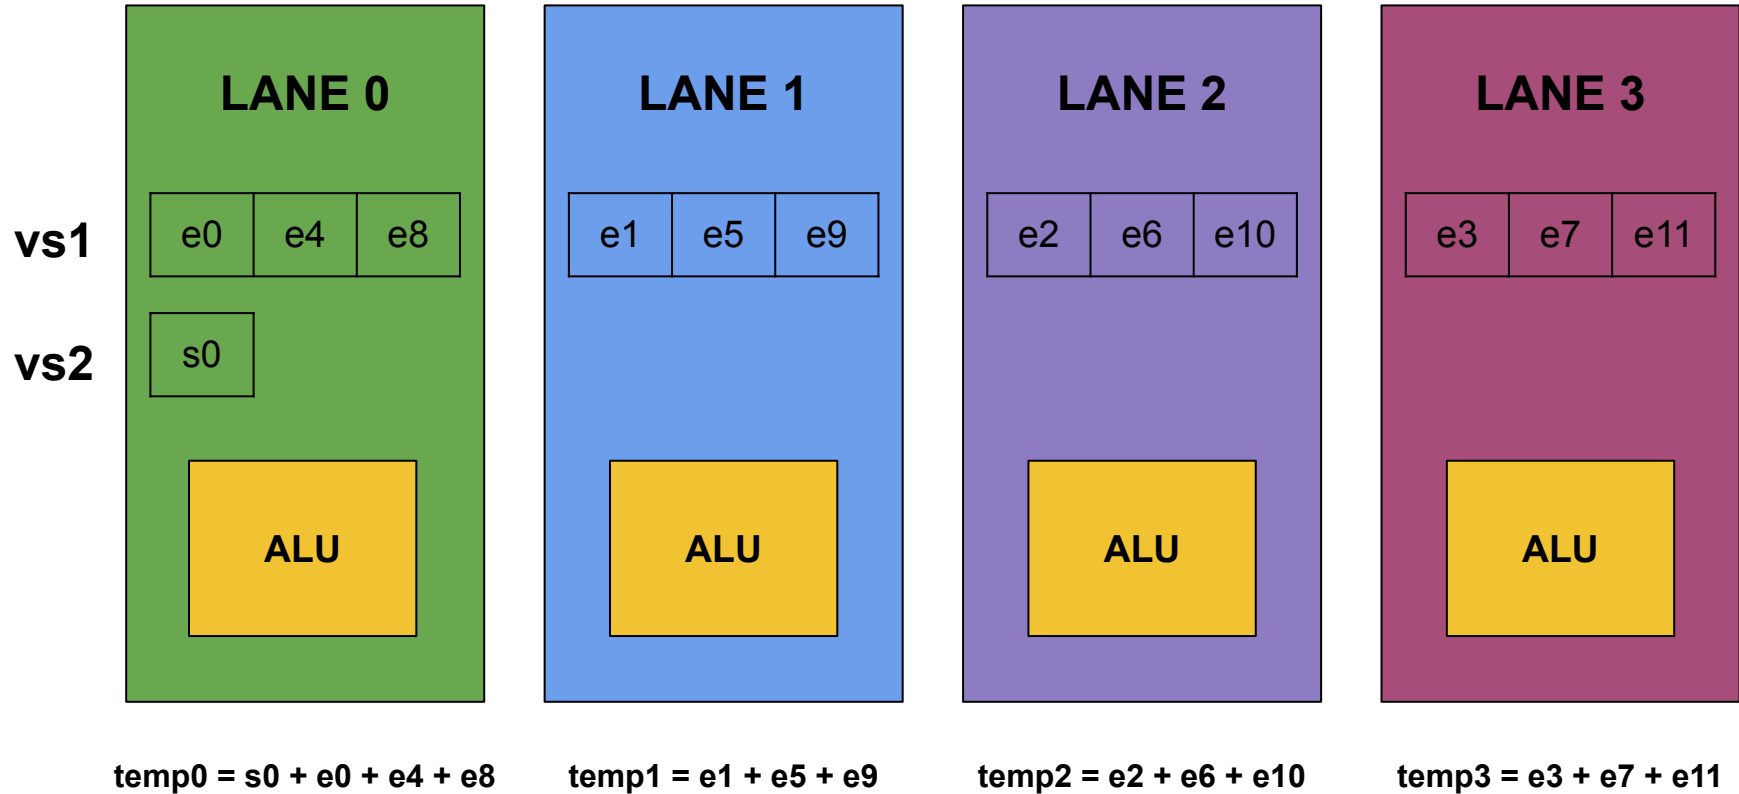
Example: vredsum vd, vs1, vs2



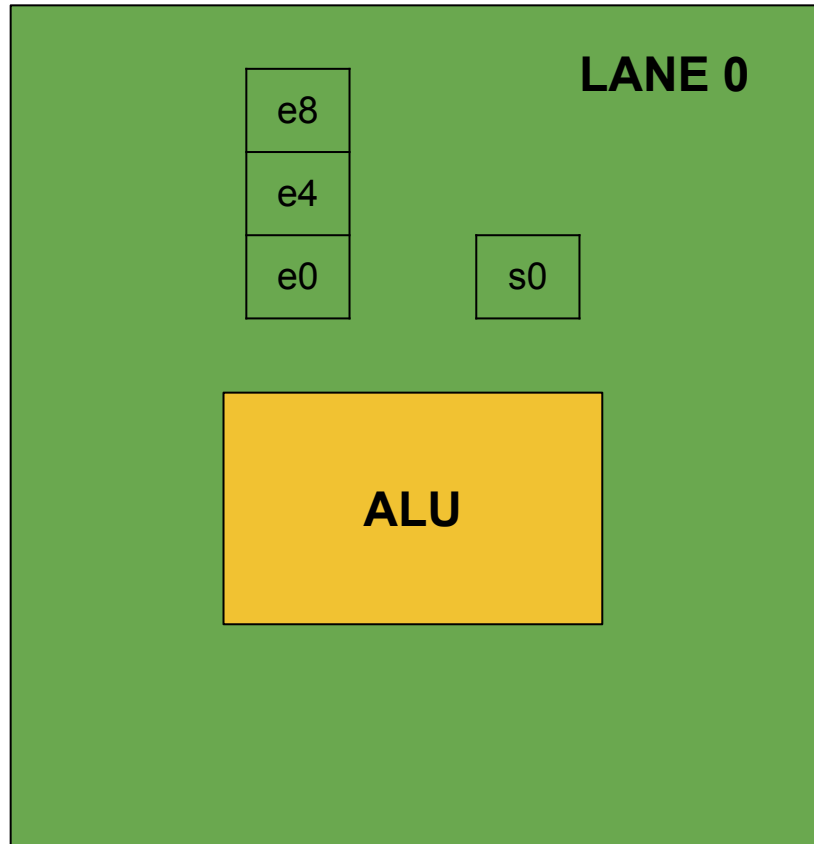
Hardware - Reductions



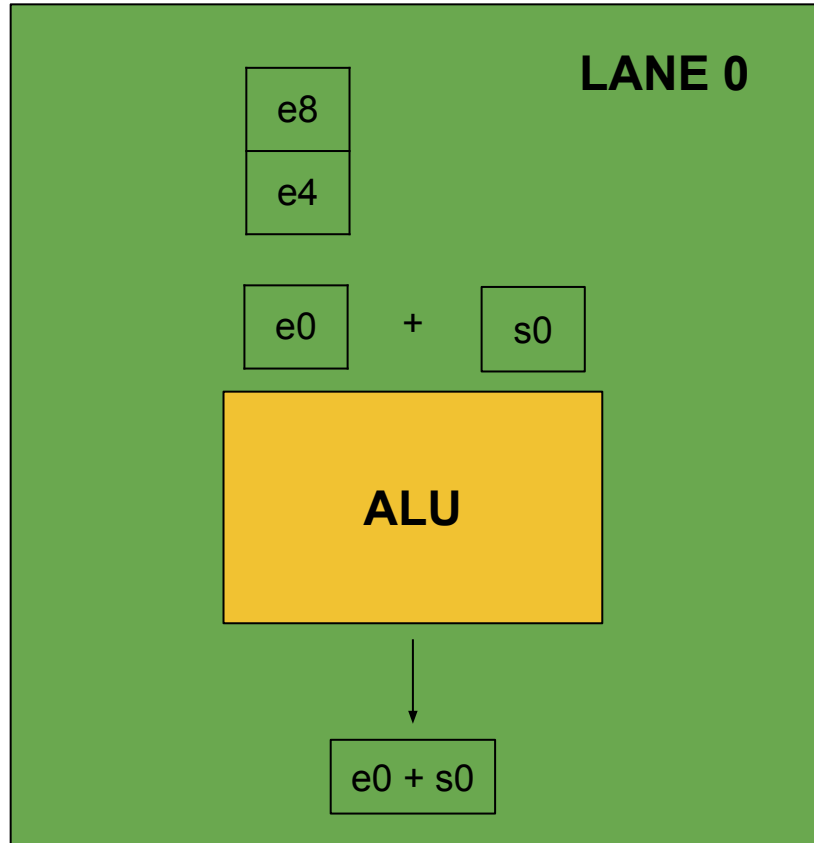
Hardware - Reductions



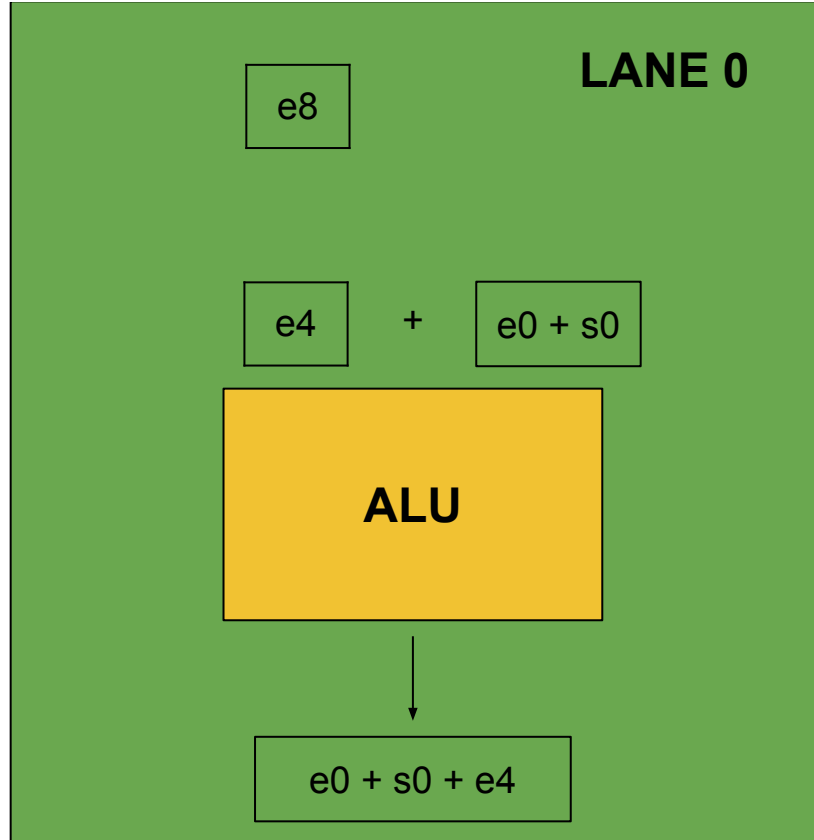
Intra-Lane reduction



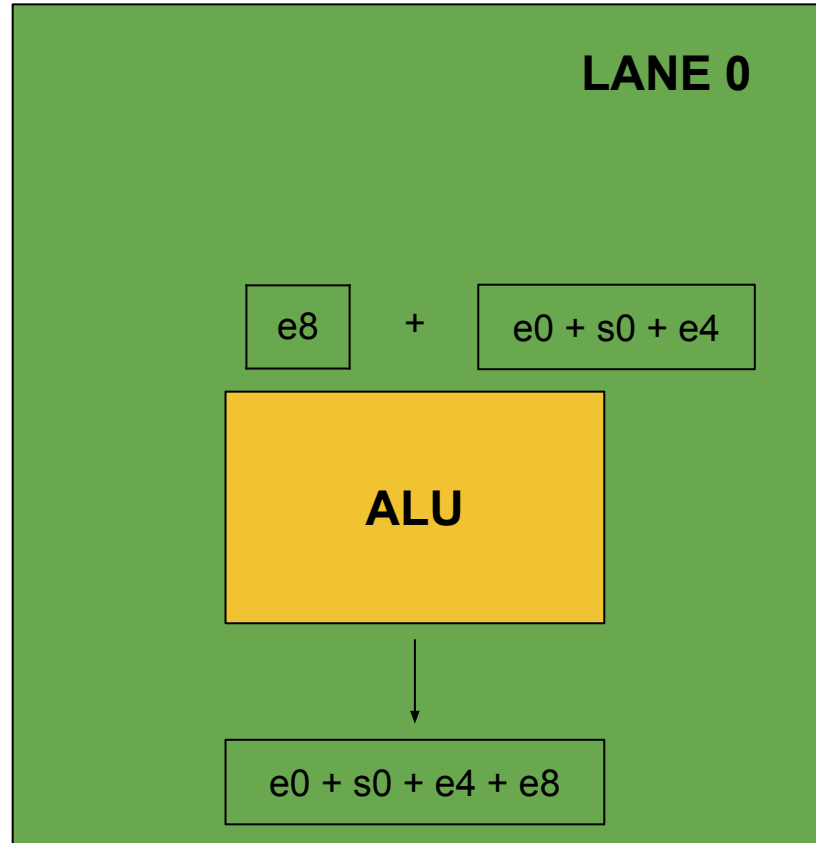
Intra-Lane reduction



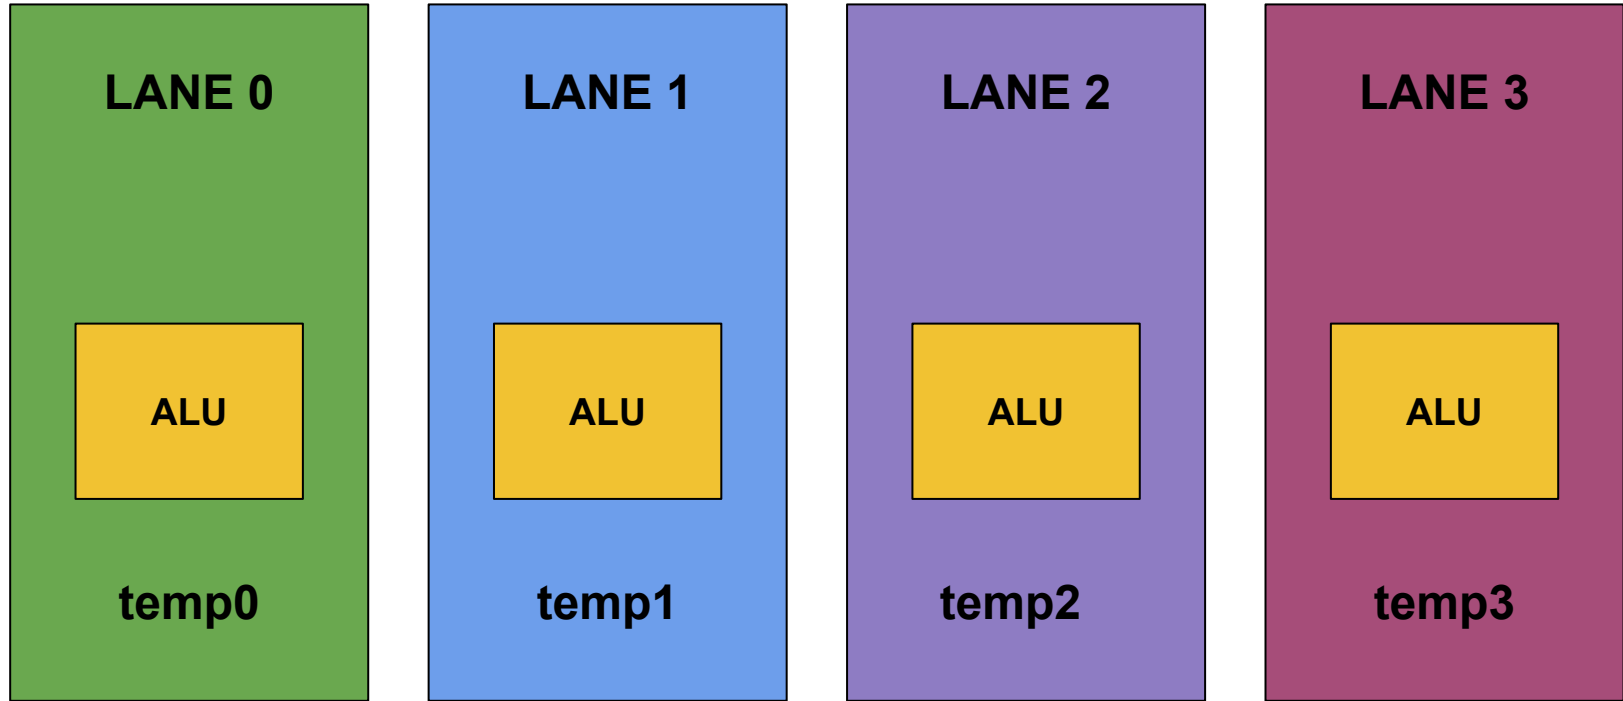
Intra-Lane reduction



Intra-Lane reduction



Inter-Lane reduction



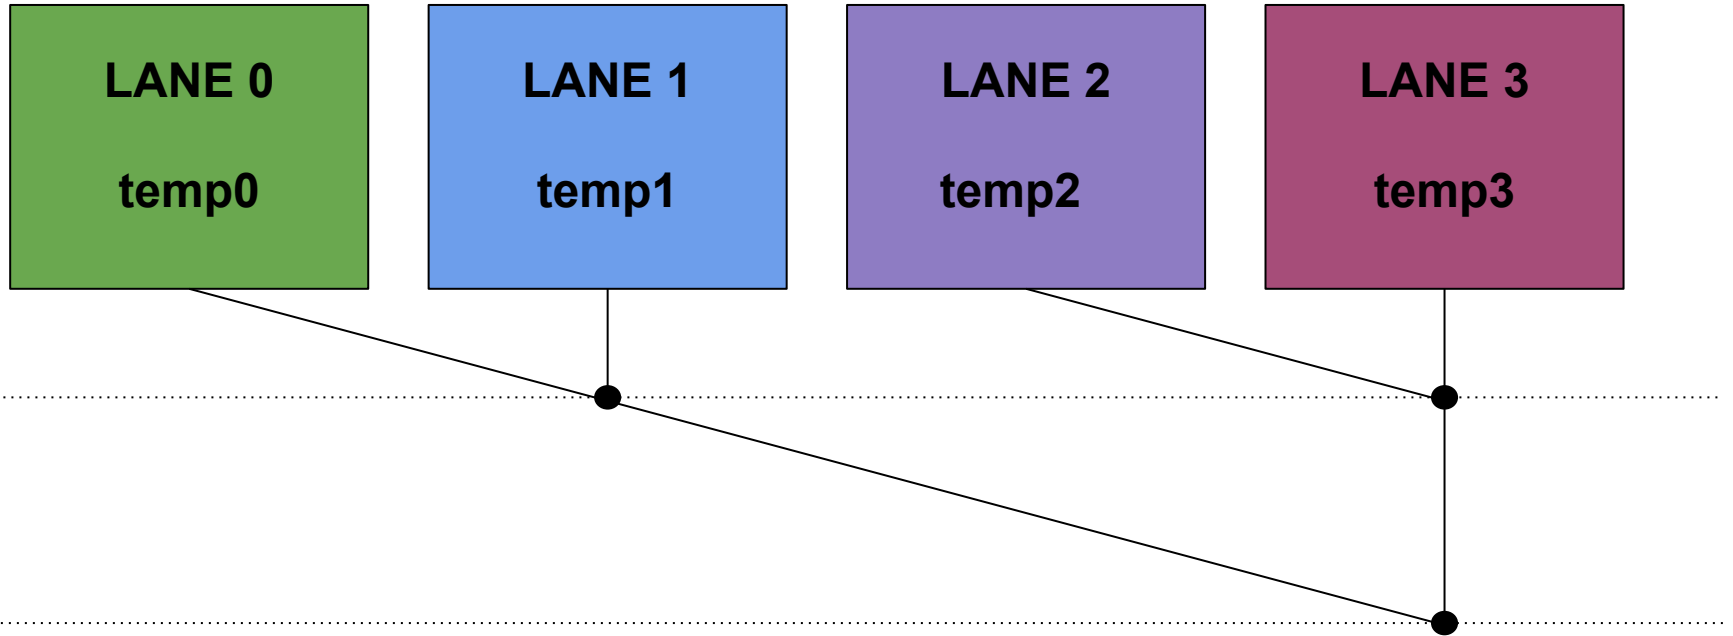
$$\text{temp0} = s0 + e0 + e4 + e8$$

$$\text{temp1} = e1 + e5 + e9$$

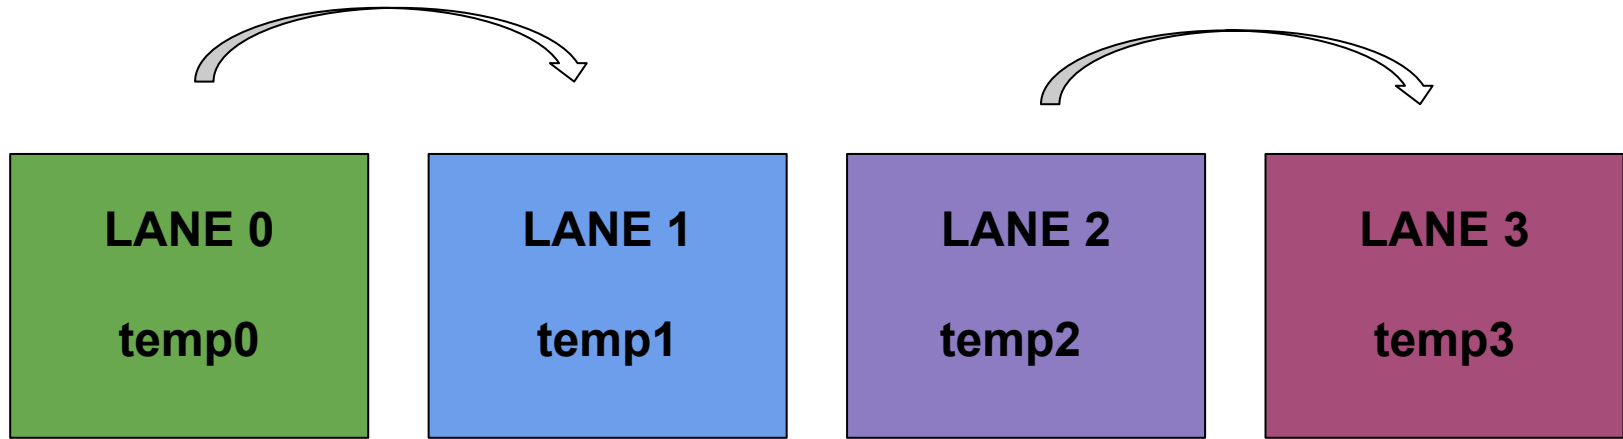
$$\text{temp2} = e2 + e6 + e10$$

$$\text{temp3} = e3 + e7 + e11$$

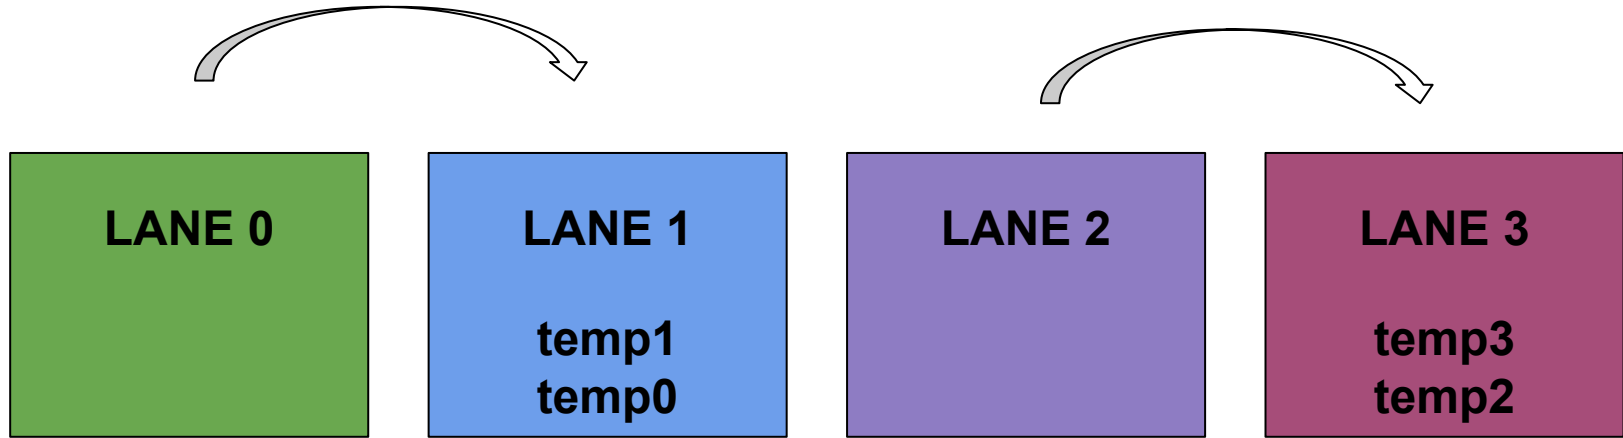
Inter-Lane reduction - Log tree





Inter-Lane reduction



Inter-Lane reduction




Inter-Lane reduction




$$\text{temp01} = \text{temp0} + \text{temp1}$$

$$\text{temp23} = \text{temp2} + \text{temp3}$$

Inter-Lane reduction

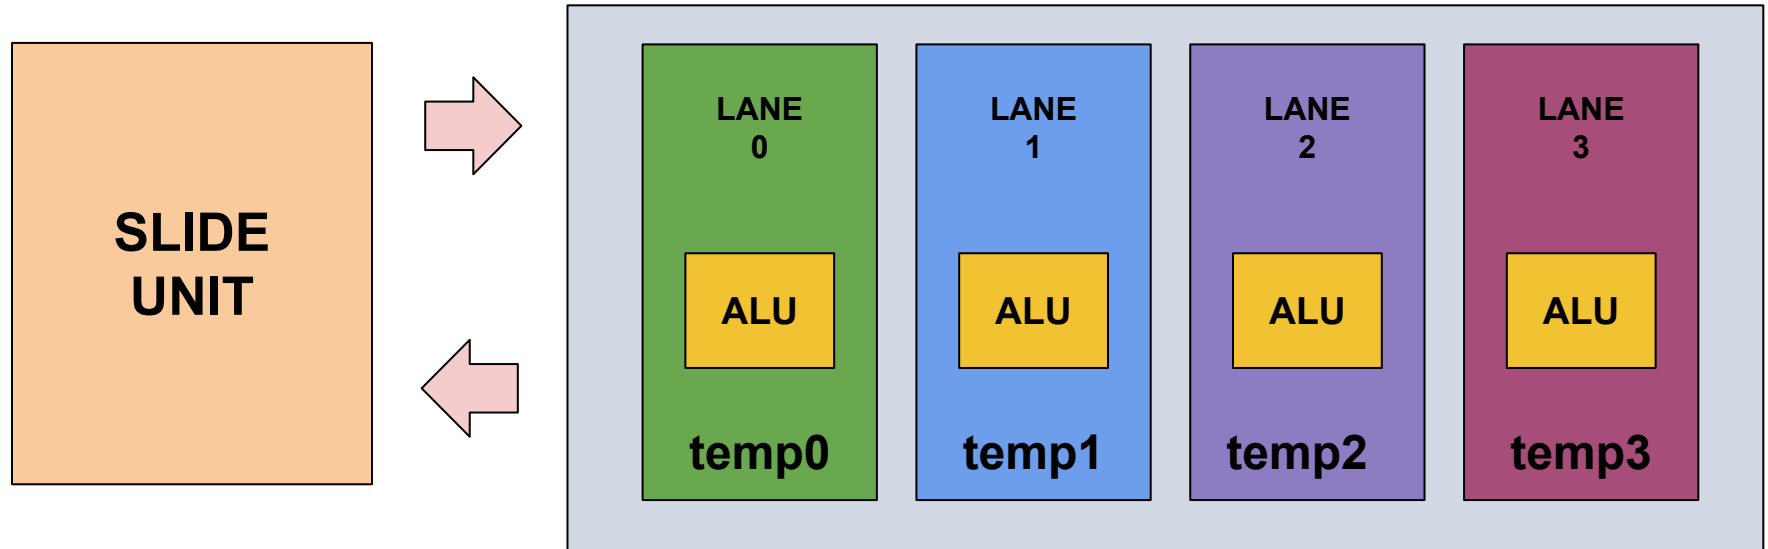


Inter-Lane reduction

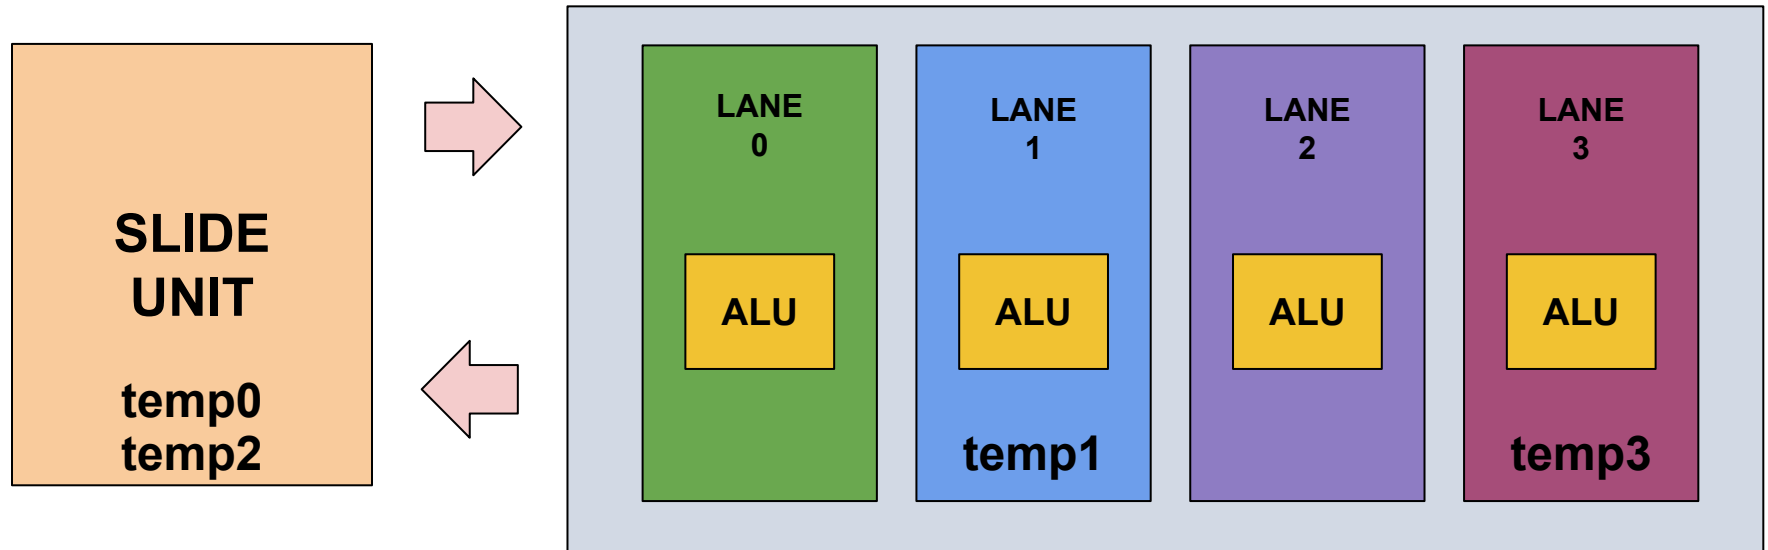


$$\text{temp0123} = \text{temp01} + \text{temp23}$$

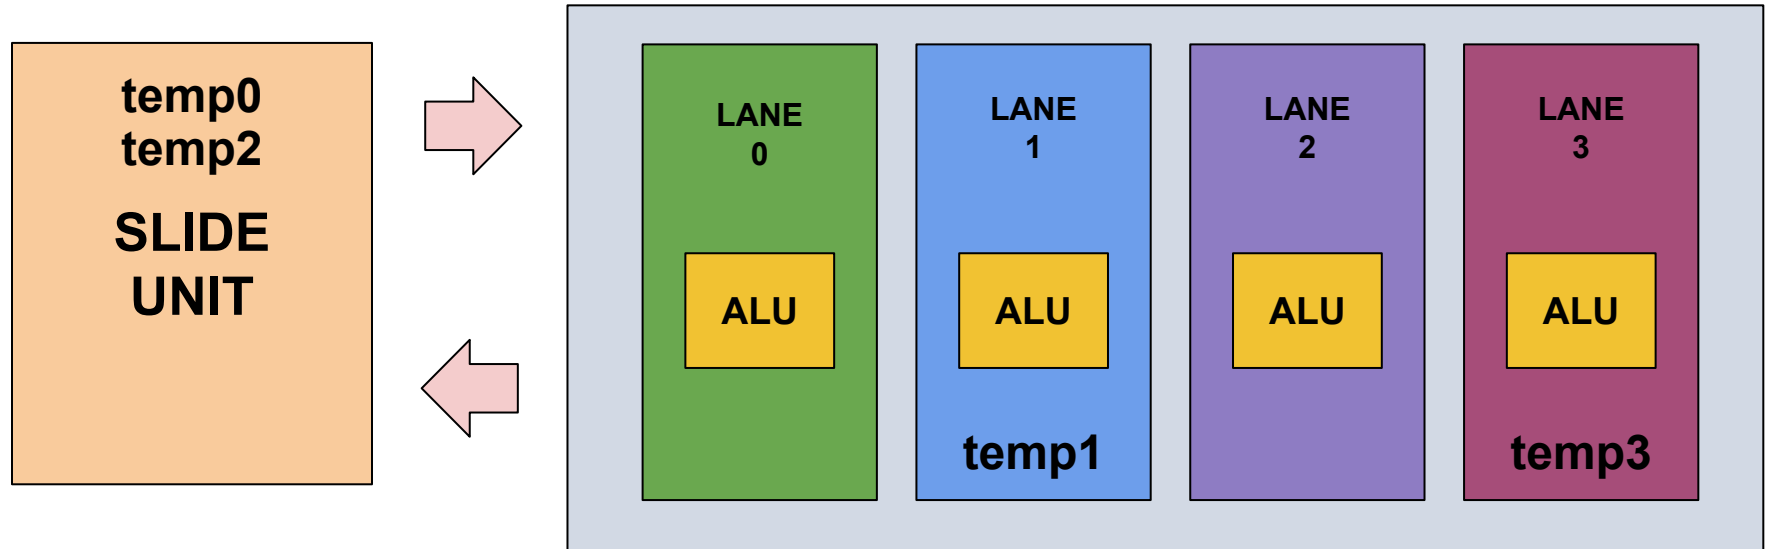
Inter-Lane reduction



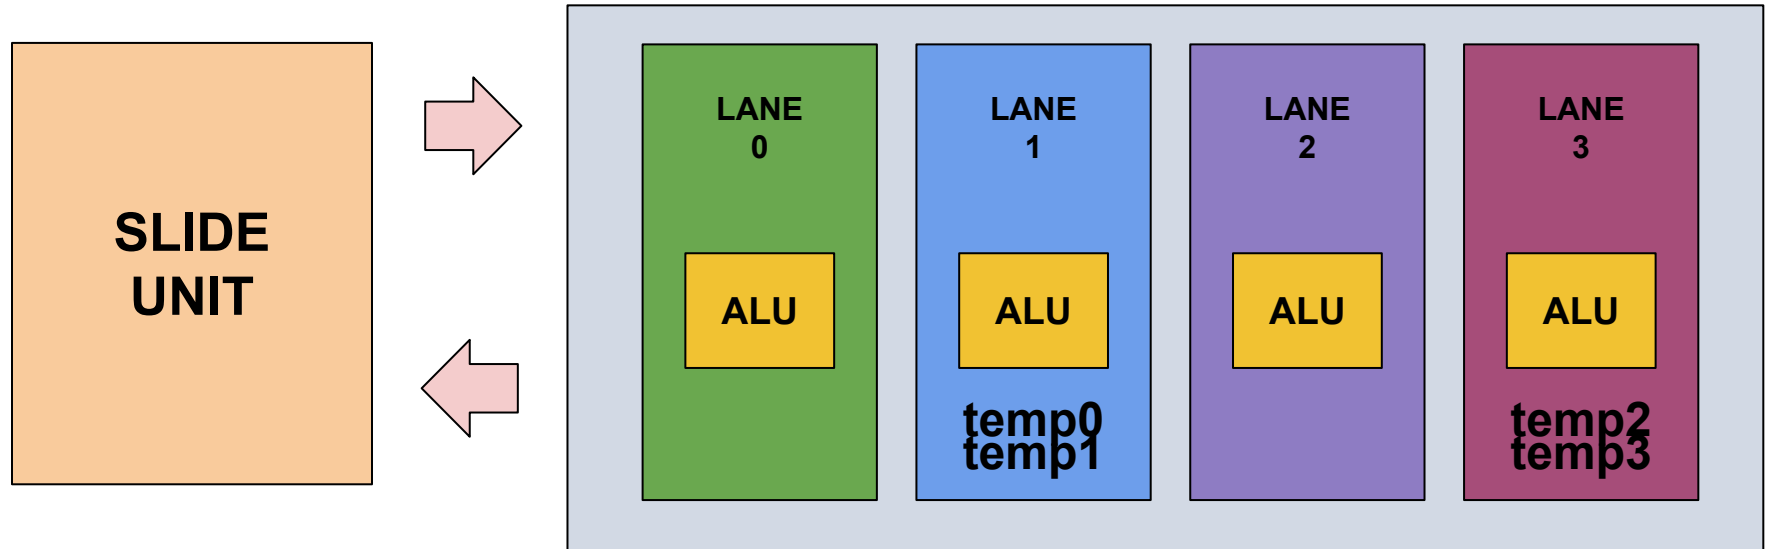
Inter-Lane reduction



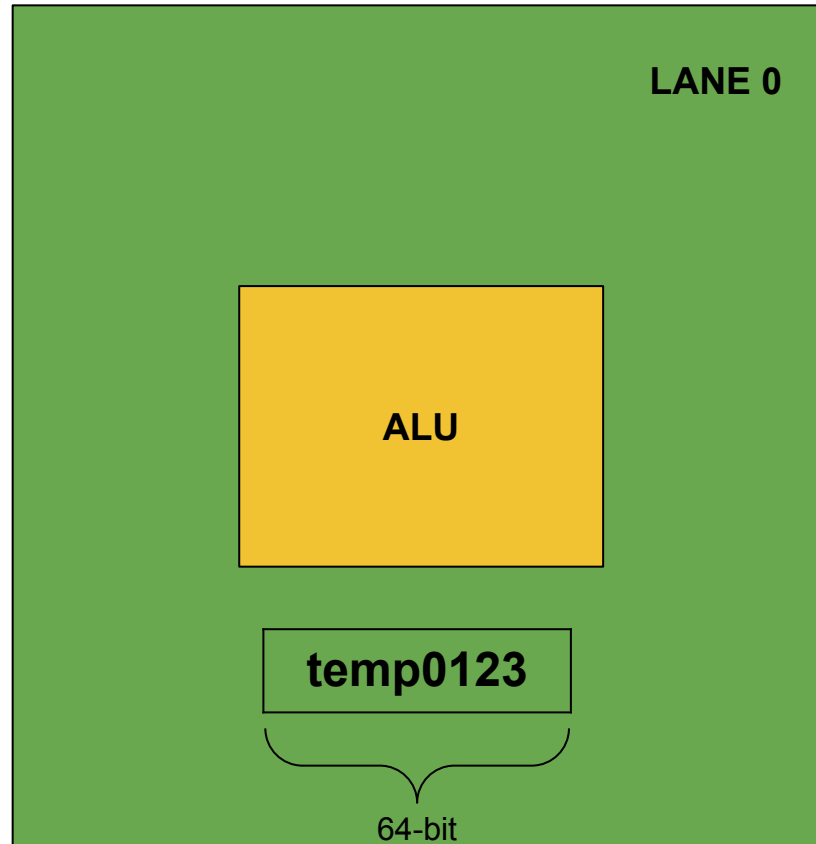
Inter-Lane reduction



Inter-Lane reduction

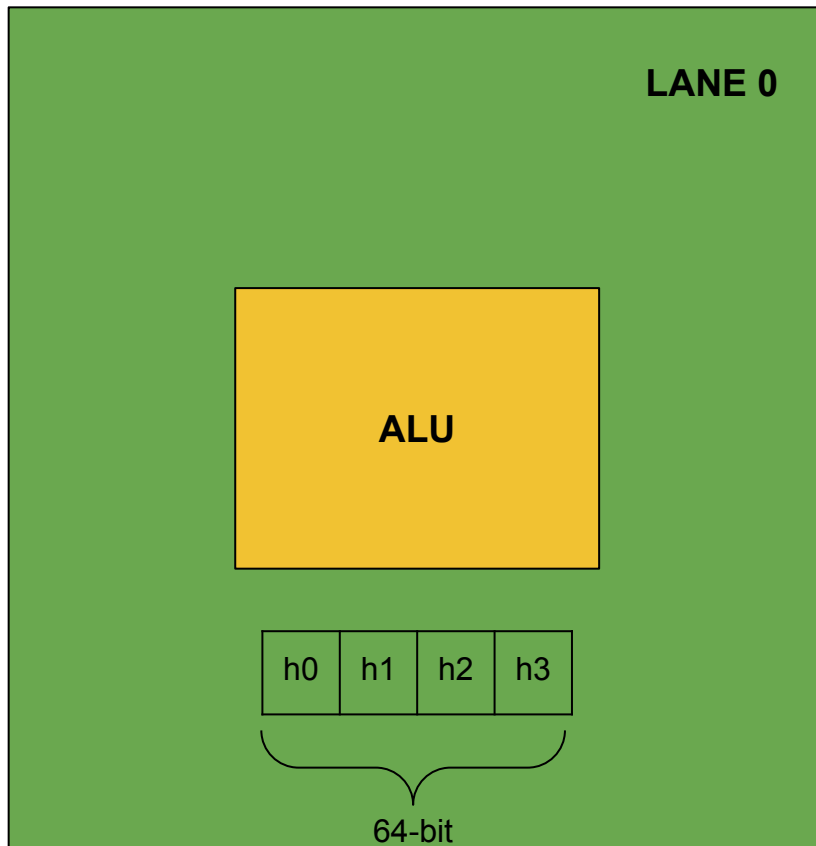


SIMD reduction




SIMD reduction (if element_width < 64)

- Reduce the 64-bit packet of N elements
- Example: four 16-bit elements
- $\log_2(N)$ operations for N sub-elements
- Example: $\log_2(4) = 2$ operations



SIMD reduction (if element_width < 64)

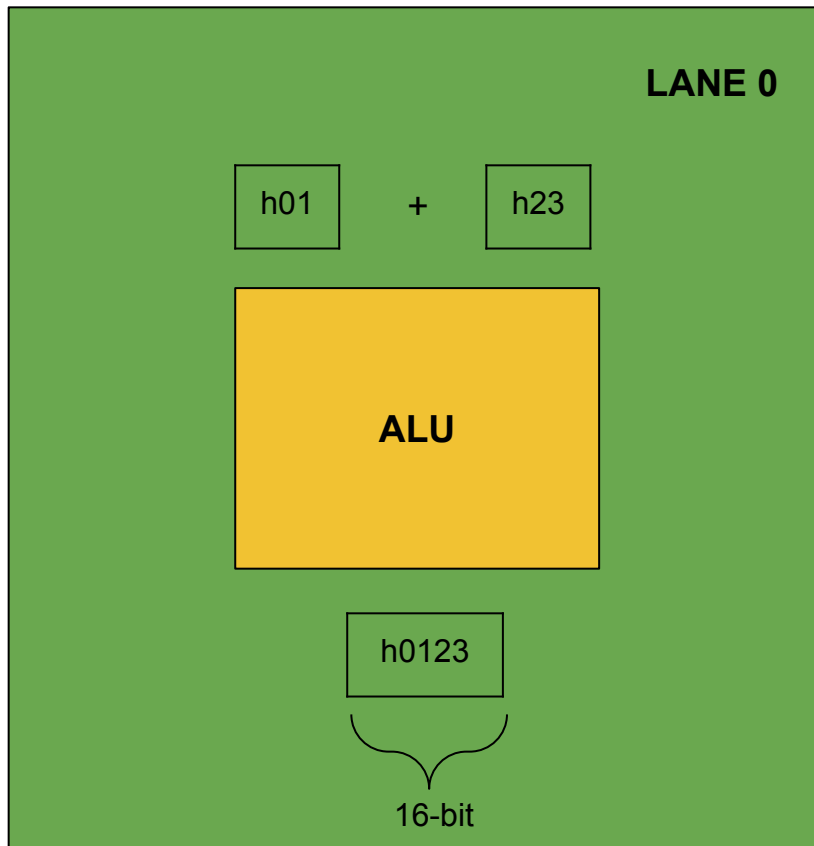
- Reduce the 64-bit packet of N elements
- Example: four 16-bit elements
- $\log_2(N)$ operations for N sub-elements
- Example: $\log_2(4) = 2$ operations



First operation

SIMD reduction (if element_width < 64)

- Reduce the 64-bit packet of N elements
- Example: four 16-bit elements
- $\log_2(N)$ operations for N sub-elements
- Example: $\log_2(4) = 2$ operations



Second operation

Reductions: evaluation

- Compare HW reductions with baseline.
- Baseline: SW sequence of slides + vector operations.

e0	e1	e2	e3	e4	e5	e6	e7	e8	e9	e10	e11
----	----	----	----	----	----	----	----	----	----	-----	-----

s0

Reductions: evaluation

- Reduce the vector with a sequence of slides and vector operations

e0	e1	e2	e3	e4	e5	e6	e7	e8	e9	e10	e11	+
----	----	----	----	----	----	----	----	----	----	-----	-----	---

e6	e7	e8	e9	e10	e11	-	-	-	-	-	-
----	----	----	----	-----	-----	---	---	---	---	---	---



Operate on half of the vector length

```
while (len > 1) {  
    len >>= 1;  
    vslidedown.vx v0, v8, len;  
    vsetvli zero, len, e64, m1;  
    vadd.vv v8, v8, v0;  
}
```

Reductions: evaluation

- Reduce a vector of 64 elements, 64-bit each
- 120 cycles for the baseline, 23 cycles for the HW implementation
- BUT:
 - This solution worsens the WNS (estimate: -3.6% frequency)
 - The baseline can be optimized:
 - Maximize throughput of the slide unit
 - Eliminate stalls with more fine-grained hazard checks on slides
 - Avoid bank conflicts
- VREDSUM baseline: good benchmark to evaluate the performance of the slide unit

SW benchmarks

- Add benchmarks to the pool and use them for verification purposes
- Optimize the benchmarks and get performance metrics

Porting to Ara already vectorized benchmarks* for RVV 0.10, with LLVM intrinsics:

- **AXPY**
- **ParticleFilter**
- **Pathfinder**
- **Jacobi2d**
- Blackscholes
- Canneal
- LavaMD2
- Others...

<https://github.com/RALC88/riscv-vectorized-benchmark-suite/tree/rvv-1.0>

from paper <https://doi.org/10.1145/3422667>

Ongoing

- Complete integer reductions (analysis of the impact on Ara clock period)
- Optimize slide unit and hazard checks
- Finish porting the already-vectorized benchmarks
- Vectorize Softmax, GaussBlur, ...
- Verification + Optimization