

Ara: Status Update

Matteo Perotti

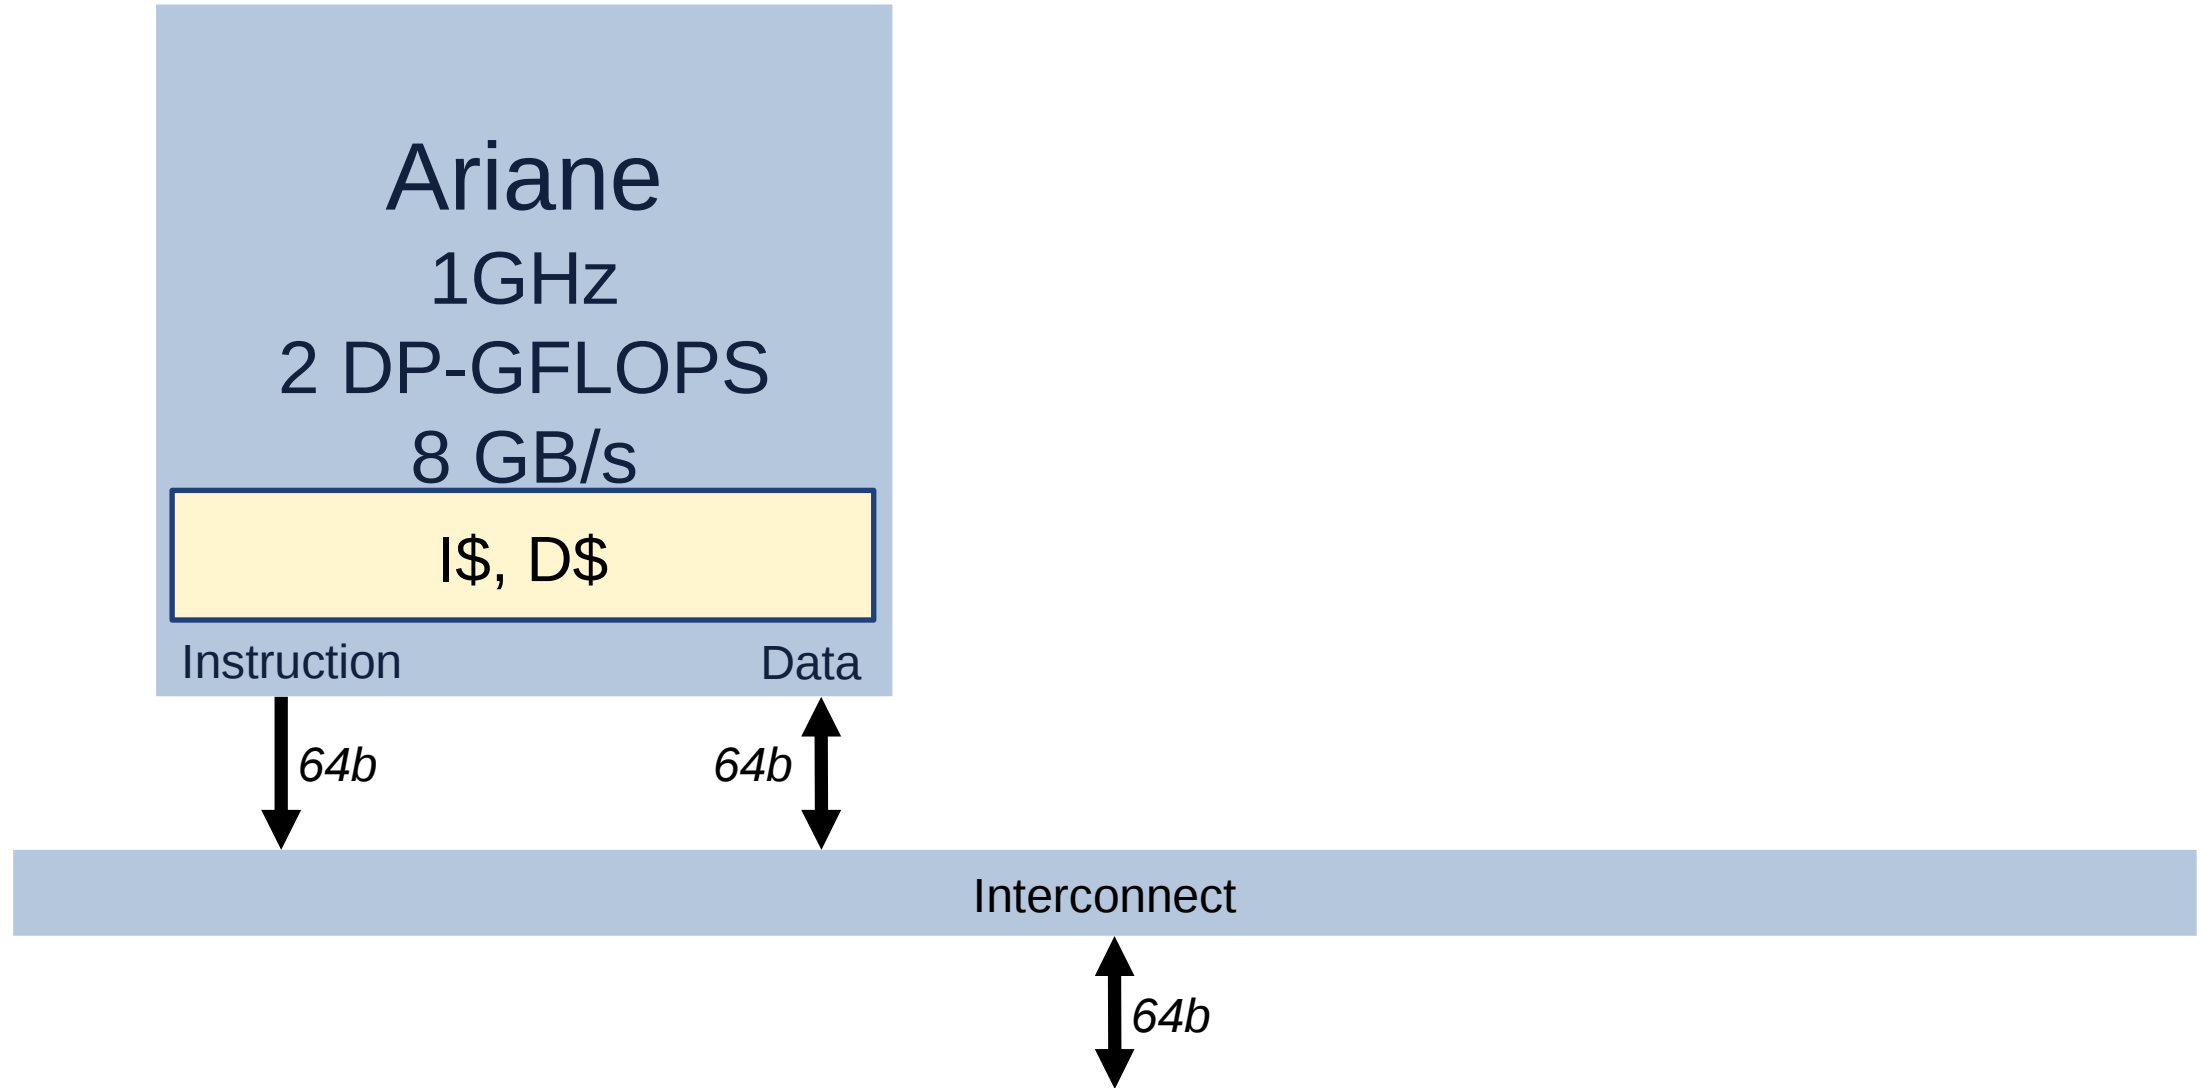
Matheus Cavalcante

PhD Students, Integrated Systems Laboratory

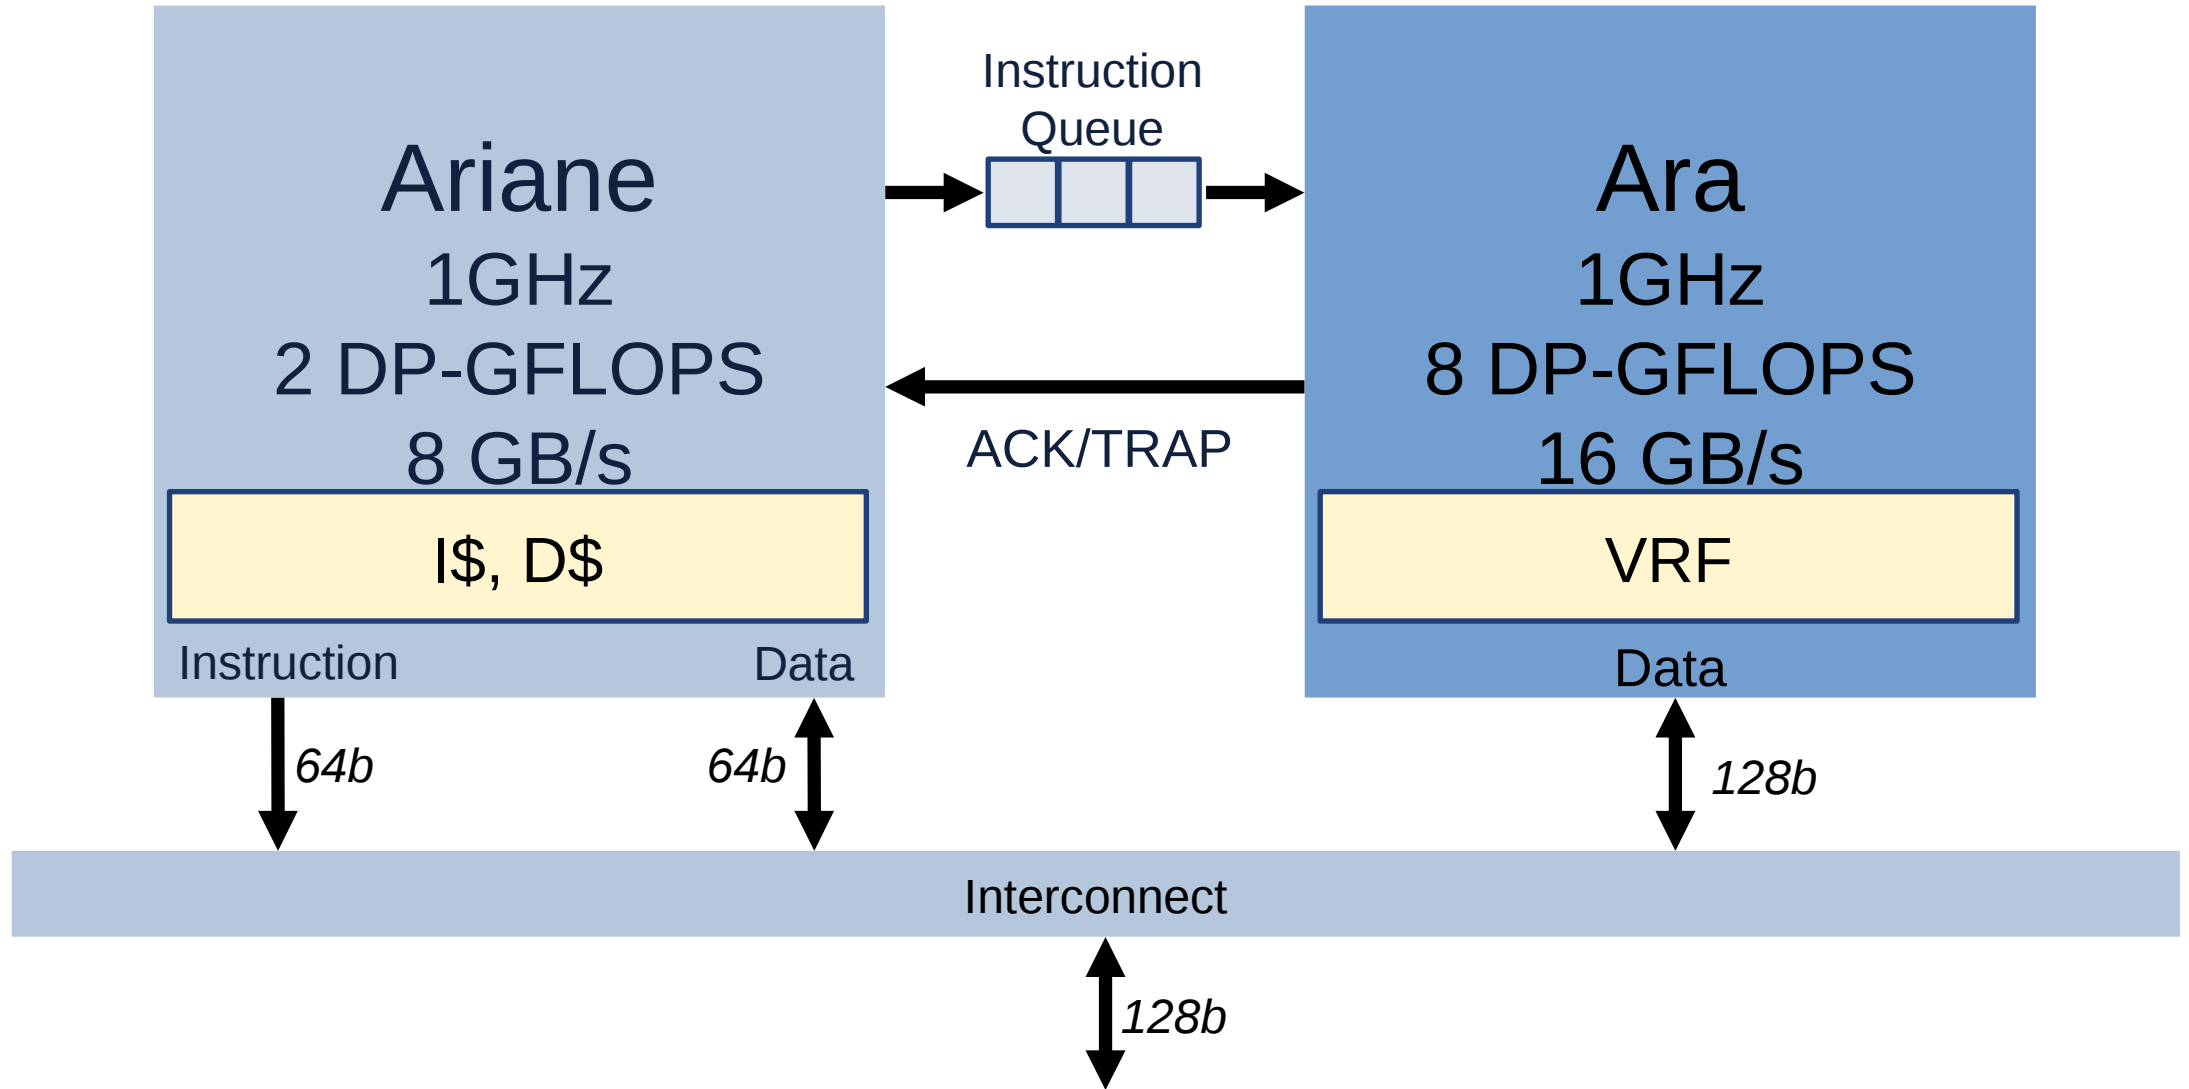
21 July 2021, Zurich



Ara: High-performance vector processor

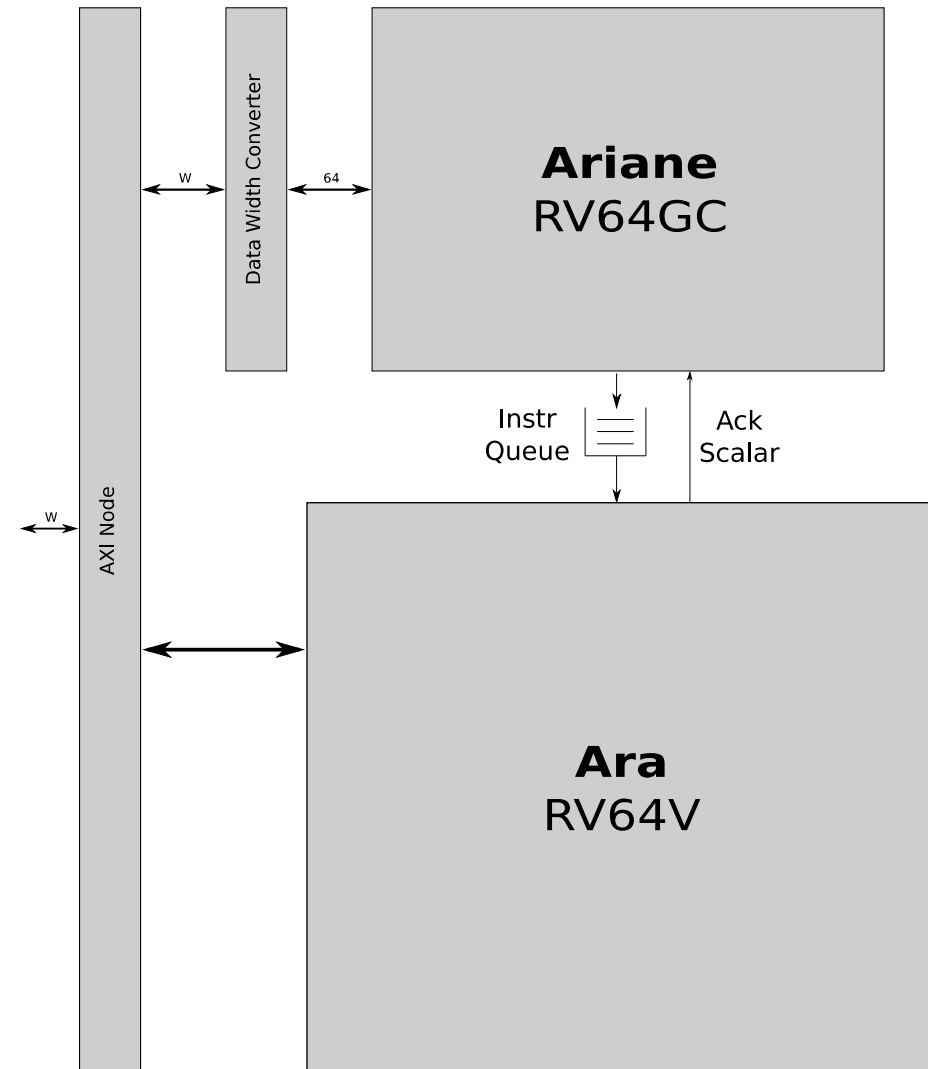


Ara: High-performance vector processor



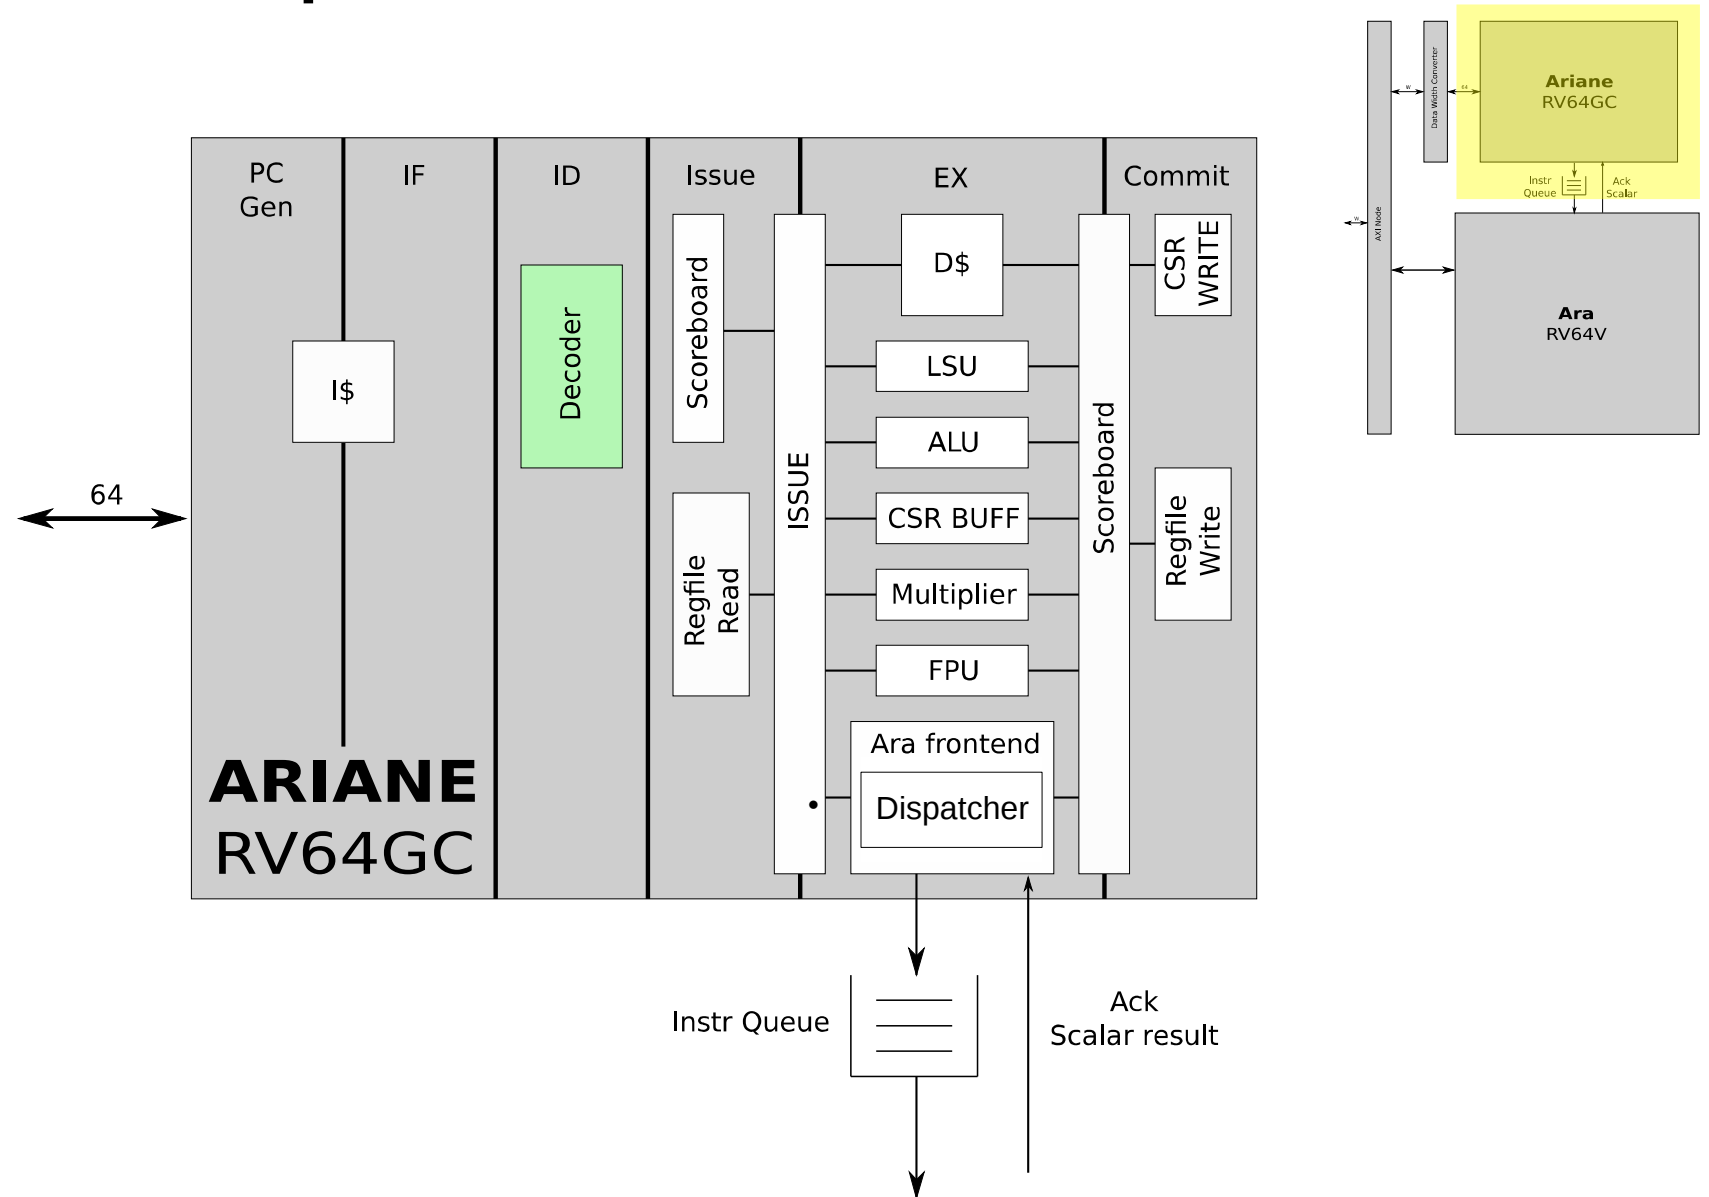
Ara - High level architecture

- Coupled with Ariane
- Ariane dispatches RVV instructions to Ara
- Ara acknowledges Ariane (Any errors? Scalar result?)
- Private memory access through AXI



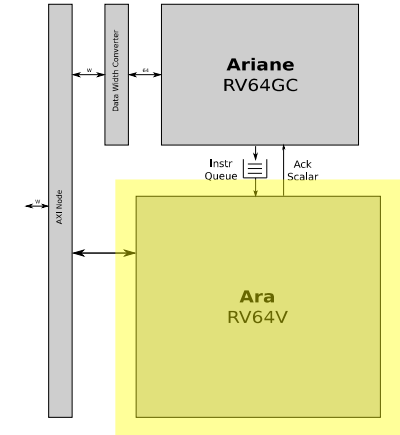
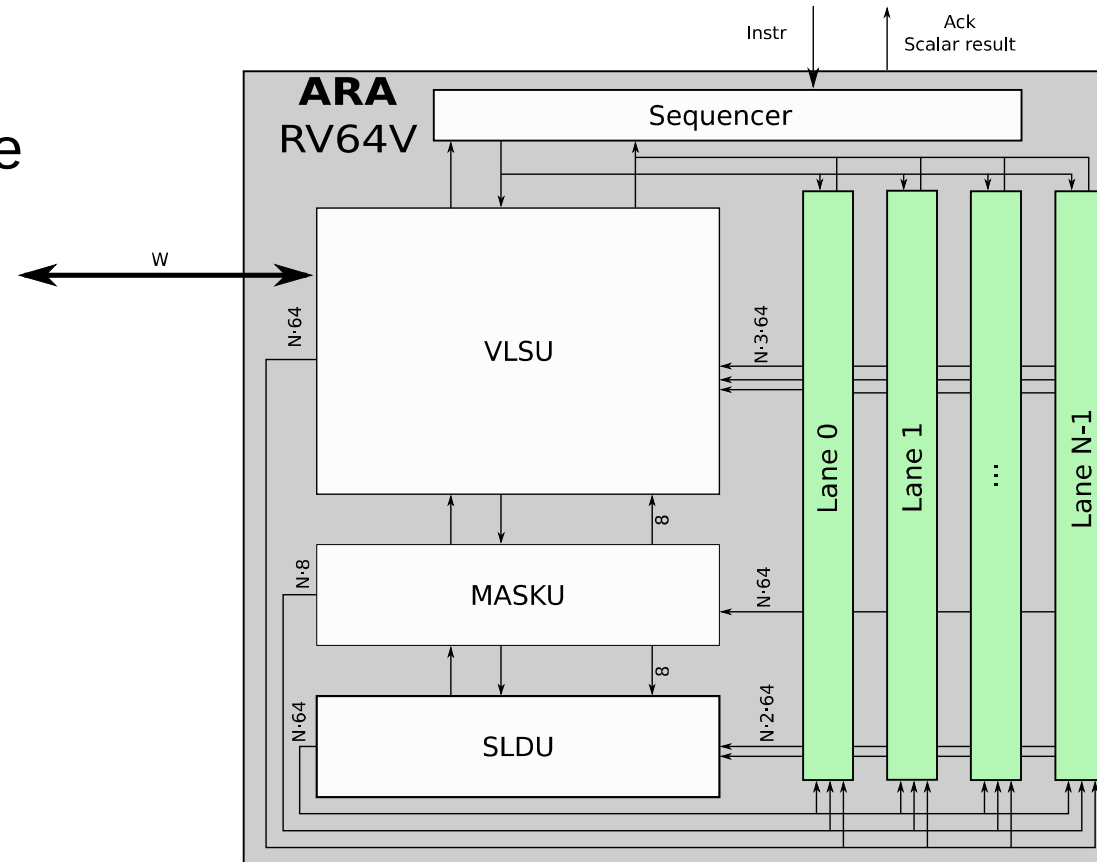
Ariane – RVV instruction dispatch

- Fetch + pre-decode in Ariane
- Scoreboard + **Issue** to Ara frontend
- Dispatch to Ara at **commit-time only**
- **Wait for Ara's answer to commit**



Ara – RVV instruction processing

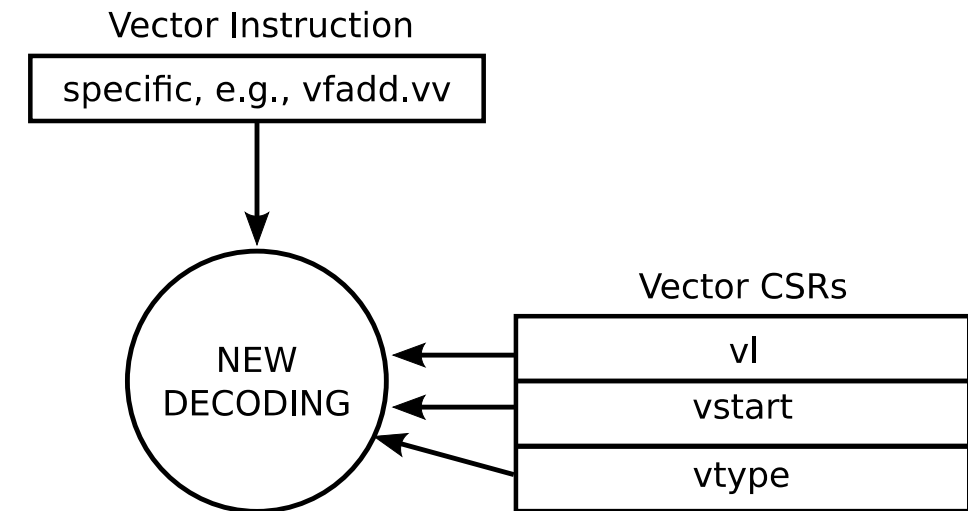
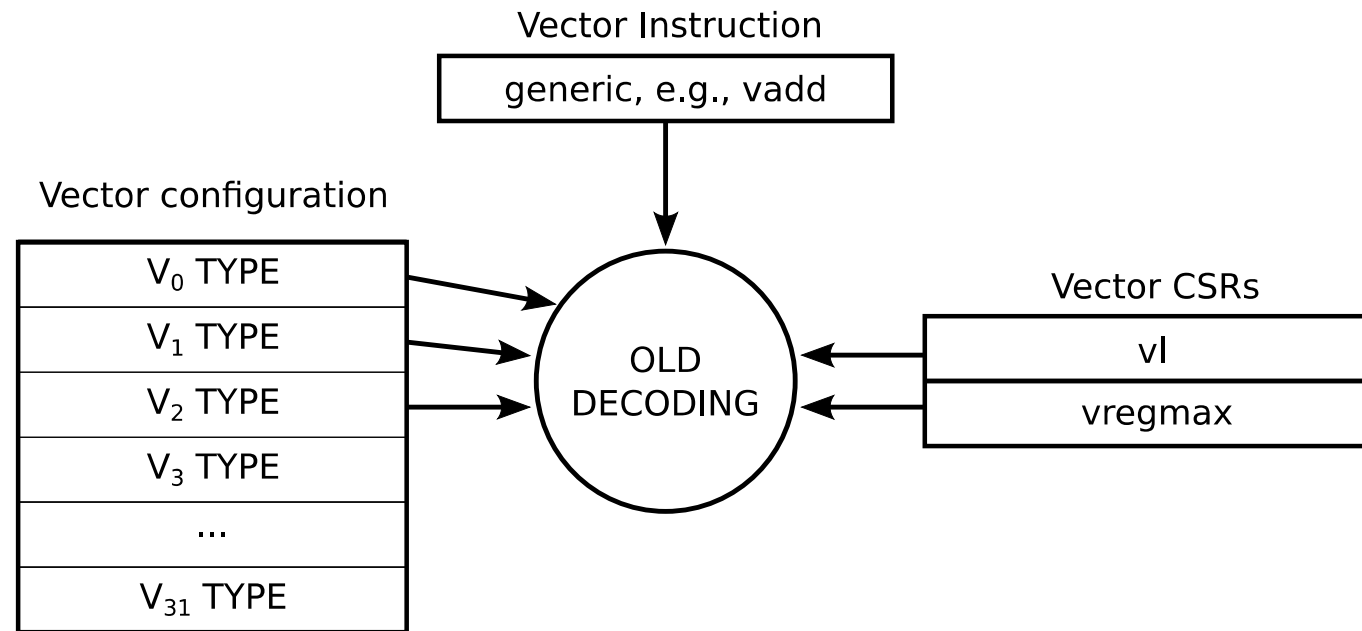
- **Sequencer**
 - Global view of Ara's state
- **VLSU**
 - **Memory (AXI) ↔ VRF** transfers
 - Address checks
- **VRF** split among the **lanes**
- **Lanes: VRF + processing**
- **New Units:**
 - **MASKU** and **SLDU**



RVV 0.5 → 0.10 changes - Encoding

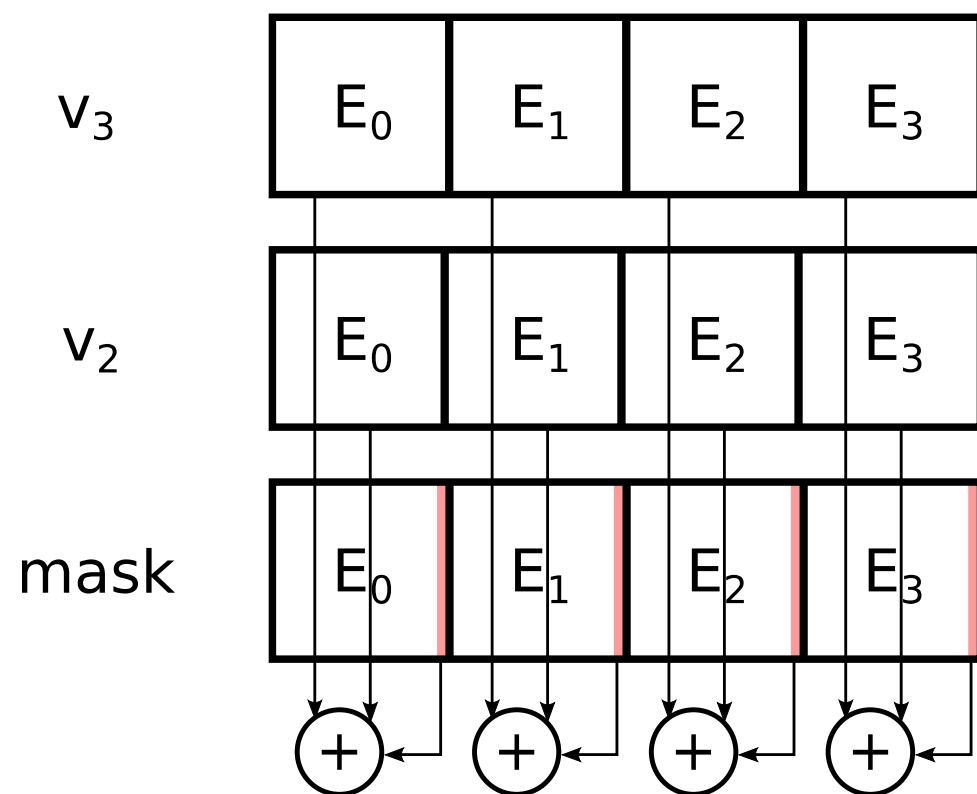
- RVV 0.5 → Polimorphic encoding

- RVV 0.10 → Monomorphic encoding

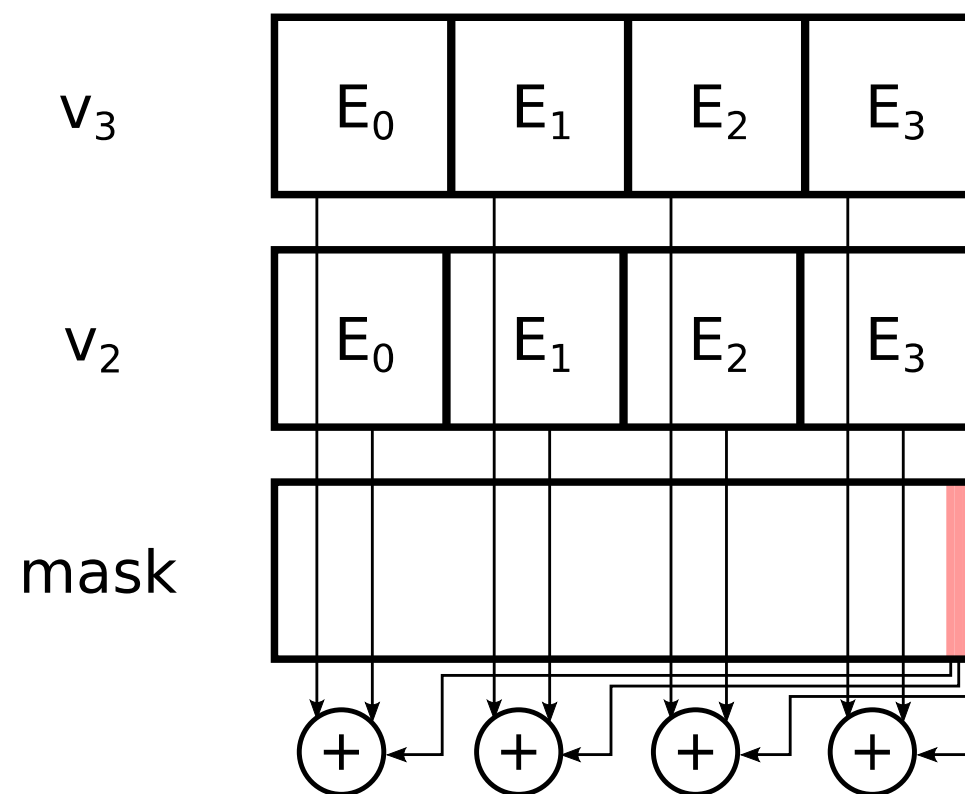


RVV 0.5 → 0.10 changes - Mask layout

OLD



NEW




RVV 0.5 → 0.10 changes - Vector length modification

Old (4 **active registers**)

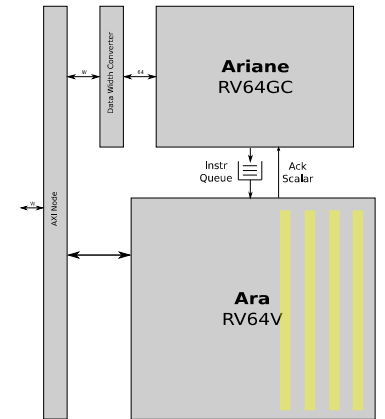
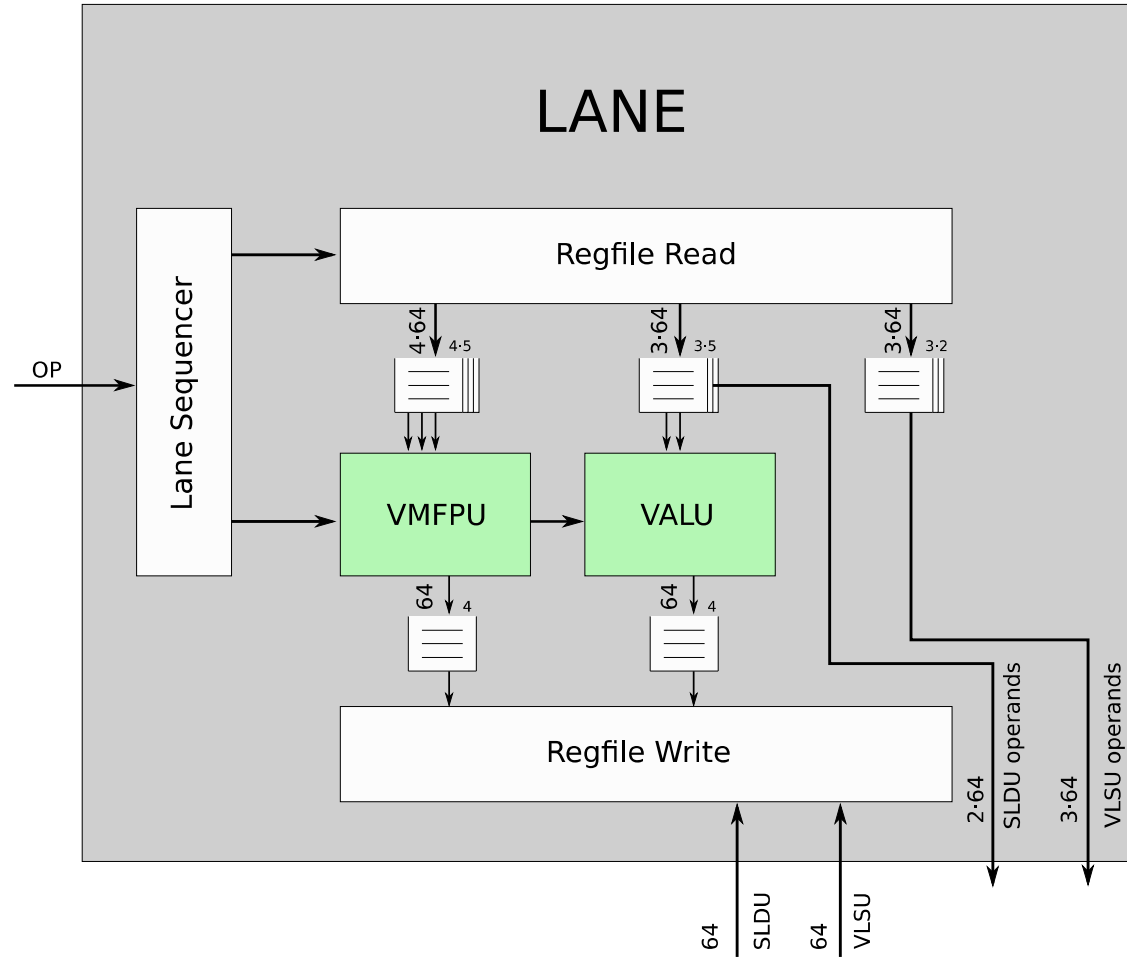


New (**LMUL 8**)



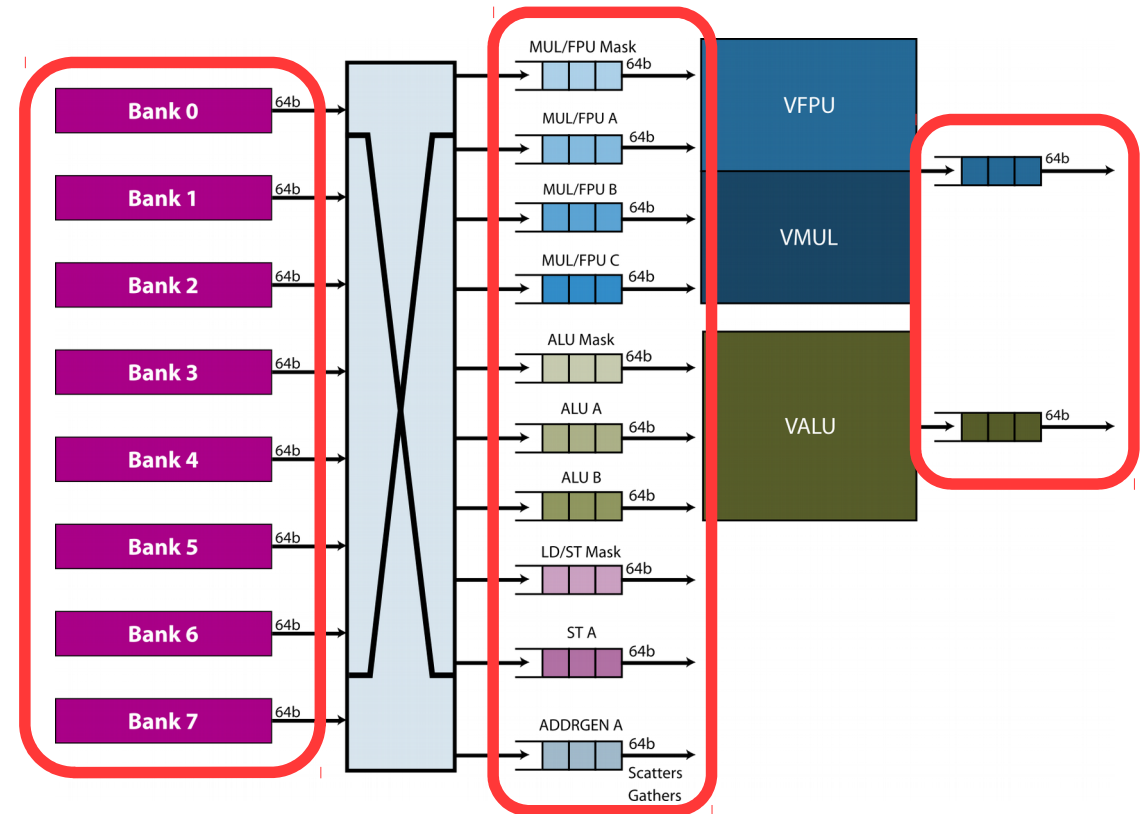
Ara Lane - RVV instruction core processing

- **Lane sequencer** tracks in-lane operations
- **8 banks per lane**
- **Multiple requests to VRF**
→ Need for an **arbiter**
- Multiple **operand queues** to feed the FUs
- **Direct write-back** from the FUs



In lane - Vector Register File, Functional Units

- **8 banks per lane** → 8 single-ported SRAM
- **Crossbar** to feed the operand queues
- **2 result queues** (VMUL and VFPU share one)
- **Arbiter** for bank conflicts



What has been done

- Memory coherence between Ara and Ariane
- **Support to RVV 0.10**, implementing:
 - **Mask Unit** (different mask layout from RVV 0.5 to 0.10)
 - **Slide Unit** (crucial for conv2d)
 - Constant-strided vector loads and stores
 - Misaligned vector memory operations
 - Narrowing and widening arithmetic and floating-point instructions
- Push **decoding** into **Ara** for increased flexibility
 - Increase decoupling between the core and Ara
 - More RV extensions-agnostic Ariane
- Yun TO



Currently ongoing

- Implementing:
 - Move whole-register
 - Whole-register memory operations
 - Mask memory operations
 - Integer reductions

Missing instructions

- Scatter-gather
- Fault-Only-First loads and stores
- Fixed-Point arithmetic
- Specific mask instructions (vpopc, vfirst, viota, vid, vmsbf, vmsif, vmsof)
- Gather/compress registers
- Floating-point reciprocal, reciprocal square root, round towards odd
- Move scalar elements from Ara to Ariane
- Floating-point reductions

SW environment – Current status

- **Compiler support**
 - LLVM + SPIKE support for RVV 0.10
 - Compiler support is not optimized
- **Refactoring our benchmark kernels**
 - `Imatmul`, `fmatmul`, `conv2d`
 - Manually vectorized
 - Intermix C + RVV assembly
 - Usage of intrinsics is possible
- **Verification**
 - Extending *riscv-tests* with ad-hoc *RVV* tests and macros

Verification

- **Support for 2, 4, 8, 16 lanes**
 - Verilator used to test all configurations
 - Check GitHub for more details